Chapter0

Table of contents

1 Motivation **2**

2 Neural Network Vision Models **2**
2.1 Convolutional-based Vision Models 3
2.2 Attention-based Vision Models 5
2.3 Heads and Transfer Learning 7

3 Image Classification **9**
3.1 Data 9
3.2 Architectures 10
3.2.1 AlexNet 10
3.2.2 VGG 11
3.2.3 ResNet 12
3.2.4 ViT 13

4 CNN-based Object Detection **14**
4.1 Object Detection Designs 14
4.2 Faster R-CNN Backbones and Heads 16
4.3 Data 16

5 Relevance to the Present Study **18**
5.1 Challenge 18
5.2 How Human Makes Judgements 19
5.3 How the Problem Will be Processed 20

6 Reference **21**

1 Motivation

Footwear examination is important to forensic investigation. The features provided by footwear prints are known as class characteristics: they may be shared by many shoes of the same make and model, but are useful for narrowing down a subject pool and providing some association between a suspect and the scene (Bodziak 2017). Collection and processing of footwear data require large amounts of human labor, hence, automating the the process of identification class characteristics is essential. With advances in neural network-based vision models, automated process has become feasible.

2 Neural Network Vision Models

Modern neural network-based vision models are mainly built on two categories of architectures: convolutional neural networks and attention-based networks. Although they use different mechanisms, these architectures are used as backbones to extract features from images. These features can be used by downstream task-specific modules, which are also referred to as heads. Task-specific heads are selected depending on the objective of the model. For example, in an image classification task, a classification head is attached to the backbone to predict image-level categories. In an object detection task, detection heads are used to predict both object classes and bounding box locations. A diagram of a neural network is shown in Figure 1.

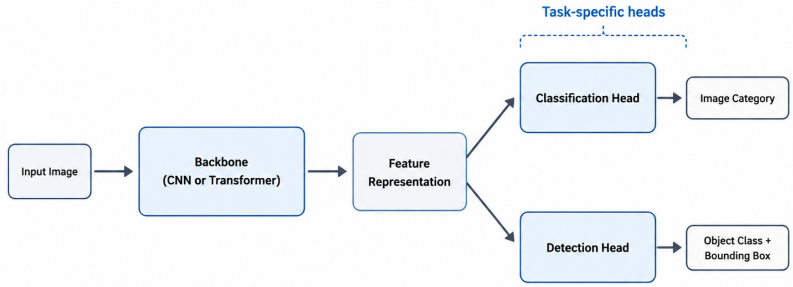


Figure 1: Vision model diagram

2.1 Convolutional-based Vision Models

Convolutional neural networks (CNNs) are one of the most widely used neural network architectures. The development of convolutional neural networks can be traced to biologically inspired models of vision.

A hierarchy model (Awang 1962) (Hubel and Wiesel 1965) proposed that the neural network in the visual cortex are structured in a hierarchical way: responses from low level cells are integrated by higher-level cells, so that more complex visual features of larger receptive field can be detected with great tolerance to positional shifts. Based on this idea, Fukushima proposed the Neocognitron (Fukushima 1980), an early neural network model for visual pattern recognition. Figure 2 shows how the Neocognitron structure corresponding to the biological hierarchy visual model. In the Neocognitron, the (U_s) and (U_c) layers provide an artificial

analogy to simple and complex cells. S-cell layers detect local visual feature, while C-cell layers combine nearby responses. A complex visual representation is built by repeating the S-cell followed by C-cell strucuter.

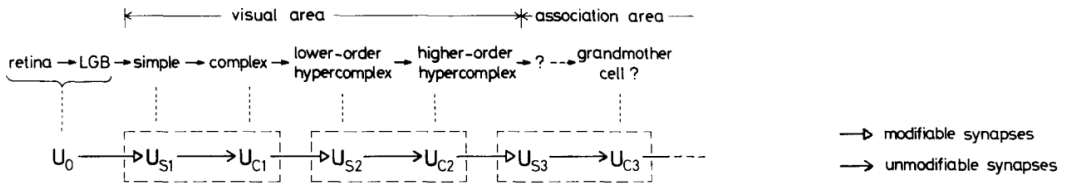


Fig. 1. Correspondence between the hierarchy model by Hubel and Wiesel, and the neural network of the neocognitron

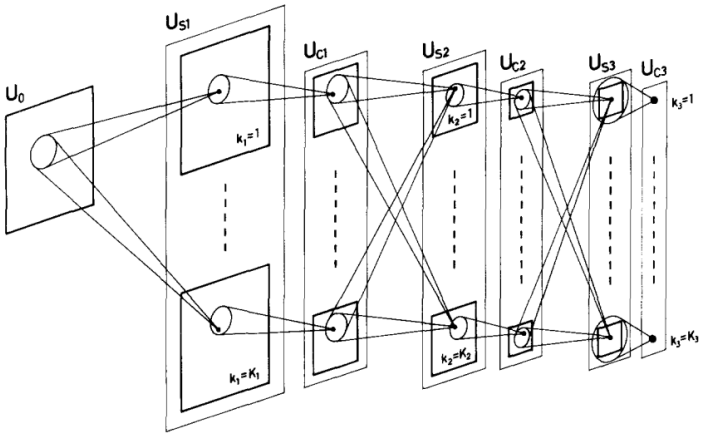


Fig. 2. Schematic diagram illustrating the interconnections between layers in the neocognitron

Figure 2: Correspondence between the hierarchy model by Hubel and Wiesel, and the neural network of the neocognitron

While the Neocognitron is considered as a biologically inspired precursor, LeNet, proposed by LeCun, represented a early form of modern convolutional neural networks. LeNet established a basic structure for later CNN architectures and enabled image features to be learned directly from image data. In this structure, convolutional filters (Figure 3) are applied to local regions of the input image to detect spatial patterns.

As a filter moves across the image, it computes responses to local visual patterns and produces

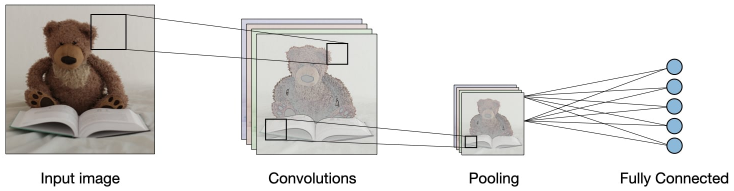


Figure 3: Convolution operation

a feature map that represents the strength of the filter response at different spatial locations. The filter plays a role similar to S-cells in the Neocognitron, in that they respond to local visual patterns and provide the basis for progressively more complex feature representations.

The use of local convolutional filters and shared weights gives CNNs strong image-specific inductive biases. Local convolutional filters implies the assumption that nearby pixels are more likely to be related than distant pixels. Shared weights allow the same filter to be reused across spatial locations, this design supports translation equivariance, in which a spatial shift in the input tends to produce a corresponding shift in the feature response. By stacking multiple convolutional layers, CNNs can gradually integrate local patterns into more complex and abstract visual representations. This local-to-global organization is conceptually related to hierarchical processing in the biological visual system.

2.2 Attention-based Vision Models

Attention-based vision models are another major class of neural network-based models. Attention mechanism is introduce in a natural language processing model named Transformer (Vaswani et al. 2017). The attention mechanism allows a model to assign different weights to different parts of the input. Heavier weights indicate relevant parts contributing more to the

learned representation, which directly models global context. This idea was later adapted to vision modeling through the Vision Transformer (Dosovitskiy et al. 2020) (ViT). The Vision Transformer inherits the attention mechanism to model region relevance. An input image is represented as a sequence of segmented visual patches, and the patches are assigned attention weights according to their relevance to the representation. In this way, the model learns relationships among different patches and captures long-range dependencies across the image. This differs from CNNs, which primarily learn local patterns through convolutional operations and gradually build broader spatial context though stacking layers.

The diagram is shown in Figure 4 The input image is first split into image patches. These patches are embedded as tokens and then processed by the self-attention module. Self-attention assigns different weights to different patches according to their relevance. The output attention maps show the spatial distribution of these attention weights across the image.

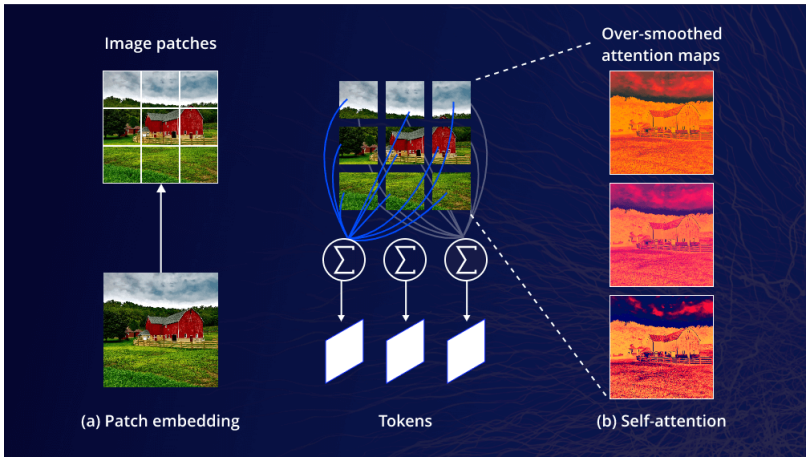


Figure 4: Vision Transformer

While CNNs assumes that local regions are more likely to be related by employing local filters,

the Vision Transformer does not impose receptive fields and models global context with a direct mechanism. Therefore, the Vision Transformer often requires large-scale training data to learn robust visual representations.

2.3 Heads and Transfer Learning

Both CNN-based and attention-based architectures can be used to extract image features. However, these features must be used by a task-specific head to complete a particular task. For image classification, a classification head maps the features to category predictions. For object detection, detection heads use the feature representation to predict both object categories and locations.

In a specific project, the task-specific head is usually selected or modified according to the objective of the study. For example, in an image classification task with seven categories, the final classification head need to be replaced by a seven-class classifier. The model is then trained by minimizing a task-specific loss function, such as a classification loss for category prediction.

In many applied vision projects, the computing resource or the number of project-specific training images are limited. Therefore, transfer learning is a commonly used strategy in the training to improve efficiency and performance. In transfer learning, a model is first pretrained on a large scale dataset, such as ImageNet (Deng et al. 2009) for image classification, and then adapted to the target dataset. Through pretraining, the backbone learns general visual features, such as edges, textures, and shapes, are considered transferable across different image

datasets and vision tasks (Yosinski et al. 2014). Then the task-specific head can be replaced or fine-tuned for the new task. The task-specific model is further trained using the target data. During the training, the pretrained backbones can be freezed or fine-tuned together with the new head, because it was found that features in deeper layers become increasingly specific to the original task (Yosinski et al. 2014). The transfer learning diagram is shown in Figure 5.

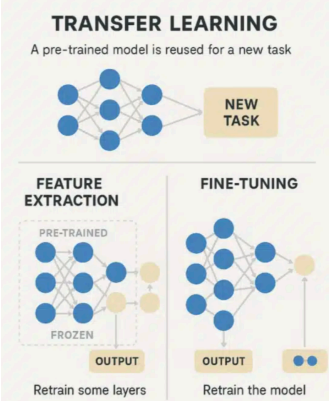


Figure 5: Transfer learning

Previous studies have demonstrated that transfer learning is effective, especially across visual network architectures and tasks (Dosovitskiy et al. 2020) (Girshick et al. 2014) (Simonyan and Zisserman 2014) (He et al. 2016).

The goal of the present study is to locate and classify footwear class characteristics. These characteristics often consist of geometric patterns that contain both local features and broader spatial context. For example, recognizing a circular pattern requires the model to identify local curvature as well as the overall closed shape. Therefore, both CNN-based and attention-based models are considered suitable for this task. Given the limited amount of labeled training data, transfer learning is employed to improve model performance.

3 Image Classification

Image classification is a fundamental task in computer vision because it provides a basic framework for learning and evaluating visual feature representations. A classification dataset consists of images and their corresponding class labels, which are assigned to the entire image. In a classification model, a classification head uses these features that output from a backbone to predict the image category. Many widely used backbones were originally developed and evaluated on image classification datasets. The representations learned through classification later became the foundation for models used in various vision tasks. This section reviews four representative architectures: AlexNet (Krizhevsky, Sutskever, and Hinton 2012), VGG (Simonyan and Zisserman 2014), ResNet (He et al. 2016), and the Vision Transformer (Dosovitskiy et al. 2020). These architectures illustrate important developments in convolutional and attention-based visual modeling.

3.1 Data

An image classification dataset consists of images and their labels. In the simplest case of single label task, each image is assigned to one category. The labeled images are divided into training, validation, and test sets. The training set is used to optimize the model parameters by minimizing a classification loss function, which measures the difference between the predicted class and observed labels. The validation set is used to identify potential overfitting, and the test set consists of held out images which are reserved for the final evaluation. During training, both the backbone and the classification head can be optimized: the backbone learns

visual feature representations, and the classification head maps these representations to class predictions.

3.2 Architectures

3.2.1 AlexNet

AlexNet, introduced by Krizhevsky (Krizhevsky, Sutskever, and Hinton 2012), was an influential convolutional neural network. It demonstrated that a large convolutional network could be successfully trained on a large-scale image dataset, and hierarchical visual features could be learned directly from image data.

The success helps to establish deep CNNs as a central approach in computer vision. Its architecture incorporated several techniques that facilitated large-scale CNN training to allow effective robust feature representation learning and large scale image classification.

Although transfer learning had been studied earlier than AlexNet, the success of AlexNet contributed to the prevalence of transfer learning. Fine-tuning a deep neural network initialized with pretrained weights, such as those from AlexNet, has been shown to outperform training the same network from random initialization (Yosinski et al. 2014). Transfer learning with networks pretrained on large-scale datasets therefore provides a general paradigm for various vision-related tasks.

Another practice popularized by AlexNet was data augmentation. Data augmentation artificially enlarges the training dataset using label-preserving transformations, which helps the

network learn more general features and reduces overfitting. AlexNet applied several augmentation techniques (Figure 6), including random cropping, horizontal reflection, and RGB intensity perturbation. These operations do not change the semantic content of the image but expose the model to different variations of the same object, improving its ability to generalize to unseen data.

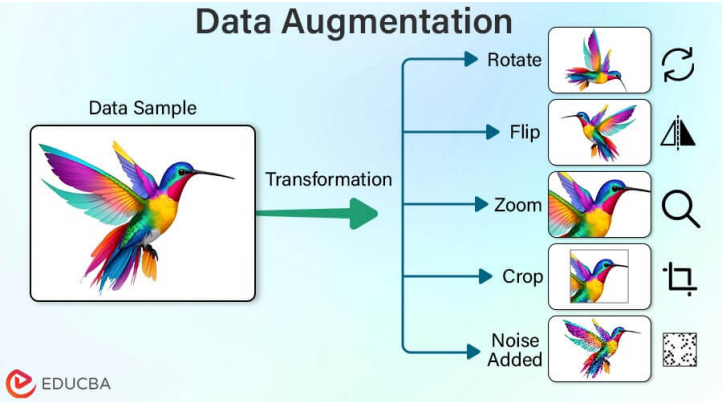


Figure 6: Data augmentation

3.2.2 VGG

Building on the success of AlexNet, deeper networks was proposed to improve feature learning, representatives of the deep networks are the configurations proposed by Simonyan et al. (Simonyan and Zisserman 2014). These configurations became widely know as VGG networks, named after the Visual Geometry Group at the University of Oxford. The two commonly known architectures are VGG-16 and VGG-19, which use a simple and uniform design based on repeated small convolutional filters (3 by 3) and contain 16 and 19 learnable layers, respectively. Compared with AlexNet, which contains eight learnable layers, VGG substantially

increased network depth, see Figure 7. The study demonstrated that this deeper architecture improved image classification performance, suggesting an improved ability to learn discriminative visual representations.

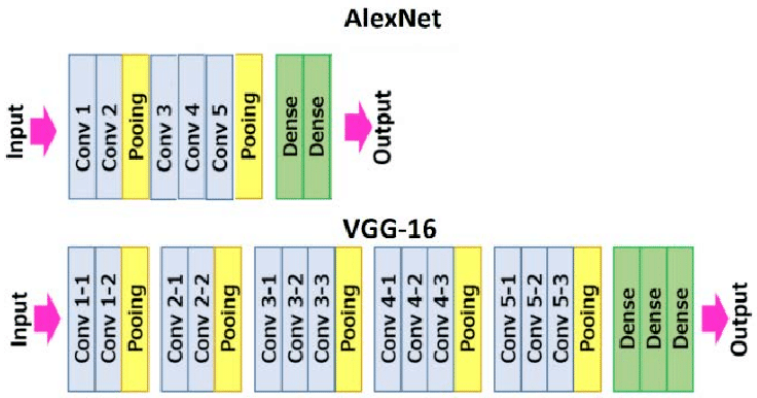


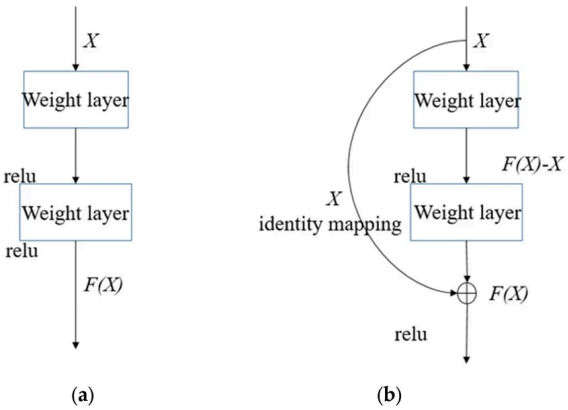
Figure 7: AlexNet: 8 learnable layers, VGG-16: 16 learnable layers

3.2.3 ResNet

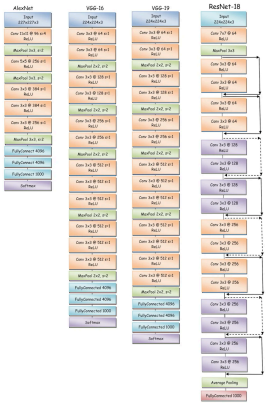
The deep layer design of CNNs brought an unexpected degradation despite better performance. It was found that adding more identical mapping layers could result in higher training error as well as test error (He and Sun 2015). This observation suggests that the performance degradation is primarily associated with optimization difficulties rather than overfitting.

ResNet architectures (He et al. 2016) addressed the problem by introducing a deep residual learning module. This module adds a shortcut connection to add the input directly to the output of a few layers. The design makes optimization easier by learning the residual modification relative to the input rather than the full mapping. If the modification is small, as in an identity mapping, the optimization becomes much easier than in the normal design (Figure 8a).

ResNet networks has five standard variants, ResNet-18, ResNet-34, ResNet-50, ResNet-101, and ResNet-152. Although ResNet-18 is the shallowest variant, it still has 18 learnable layers, which is deeper than the VGG-16 (Figure 8b). ResNet achieved improved image classification performance while enabling much deeper networks to be trained effectively.



(a) Residual block and normal block



(b) Learnable depth: AlexNet (8 layers), VGG-16 (16 layers), and ResNet-18 (18 layers)

3.2.4 ViT

While CNNs rely primarily on local convolutional operations to extract visual features, the Vision Transforner (ViT (Dosovitskiy et al. 2020)) introduced a different approach that captures global dependencies across an image. Attention weights are learned to represent relationships among image patches. Unlike CNNs, which have inductive biases such as locality and translation equivariance built into the architectures, ViT uses global self-attention and imposes relatively little two-dimensional neighborhood structure. Therefore, ViT performance can be more sensitive to the amount of project-specific data

These architectures were initially proposed as classification models, which includes feature

extraction and class prediction. The major architectural difference lie in the feature extraction designs (backbone) rather than the classification heads. The classification heads serve a similar role to map extracted features to class predictions.

Since backbone reflects the main architectual characteristics of models designed for the same task, architecture names are commonly used to describe the unique underlying feature extraction design when these models are adapted as backbones for other tasks. For example, the same detection framework can be implemented with different backbones, such as VGG or ResNet, and these variants are often referred to by their backbone architectures.

!!!!!!!!!!!!!!!!!!!!!!!!!!!! Representative backbones options for visual feature learning, and it is therefore worthwhile to evaluate their performance in classifying footwear characteristics. !!!!!!!!!!!!!!!!!!!!!!!!

4 CNN-based Object Detection

To achieve the goal of the present study, footwear characteristics must be both localized and classified to automate the process. Therefore, the task is formulated as an object detection problem rather than an image classification problem.

4.1 Object Detection Designs

In the detection workflow, the input image is first processed by a visual backbone to extract feature representations. A detection head then uses these representations to locate and classify

target patterns. Object detection models can be broadly categorized into two-stage and one-stage models. Two-stage detectors first generate candidate regions that may contain objects and then classify and refine these regions, multiple networks are trained at the same time. Faster R-CNN is a representative two-stage detector (Figure 9 upper panel). In contrast, one-stage detectors predict object locations and class labels directly from image features without a separate region proposal stage (Figure 9 lower panel). YOLO models are representative examples of this approach.

One-stage detectors simplified the architecture by excluding the additional region proposal stage, this design reduces computational cost and enables faster inference, making them suitable for real-time applications. However, without a separate proposal and refinement process, one-stage detectors may have relatively lower localization accuracy compared with two-stage detectors. The present study does not require real-time detection and prioritizes accurate localization, Faster R-CNN was selected as the detector.

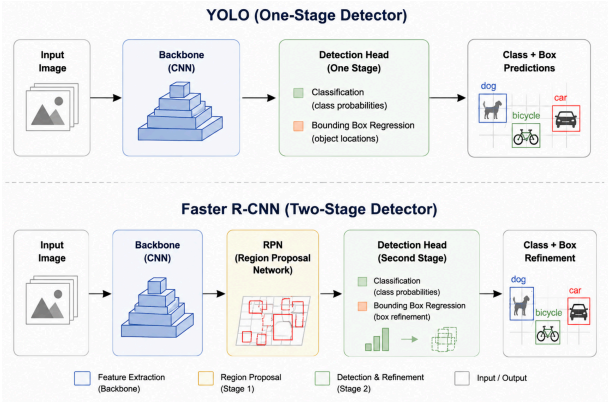


Figure 9: One stage YOLO and two stage FasterRCNN

4.2 Faster R-CNN Backbones and Heads

Faster R-CNN framework consisting of a backbone, a Region Proposal Network (RPN), and a second-stage detection head. The backbone is responsible for converting the input image into feature representations that retain visual and spatial information in the image. Based on the information, the RPN generates candidate regions that possibly contain the objects. The candidate regions will be send into the detection head. The detection head has two parallel branches: a classifier that predicts categories for the regions, and a regressor that adjusts the coordinates of the proposed box.

Faster R-CNN implementations can be realized by changing the backbone while the overall framework remains unchanged. As mentioned previously, backbone designs fundamentally affect the feature representations provided to the detector, and consequently, detection performance. Different backbone designs can provide different levels of semantic abstraction and retained spatial information (Lin et al. 2017). The retained spatial resolution is particularly important for localization because it determines how much positional detail is preserved.

4.3 Data

Object detection has a different data structure than the classification task. To obtain the ability of predicting spatial location, the training data must provide positional information in addition to the class label. Input images are paired with these human-provided annotations, which serve as the ground truth that guides model training and evaluation. For example, in

Figure 10, the image is annotated with both bounding boxes to indicate spatial locations as well as class labels for objects in the bounding boxes. An individual object detection image allows multiple target instances, each instance associates with a class label and a spatial annotation. The annotated regions can belong to the different classes, and it is common to find them overlap.

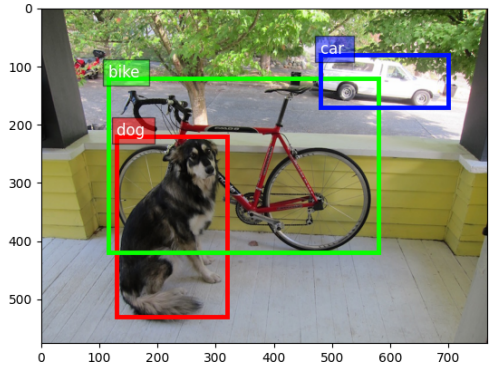


Figure 10: Object detection data

Low quality input would limit the network performance, hence, annotation quality is an important consideration in object detection. Previous studies have shown that small discrepancies in location annotation, such as different box sizes, and noisy can influence localization performance, especially for small objects (Murrugarra-Llerena, Kirsten, and Jung 2022) because its susceptibility to errors. In addition, deviation from intended localization has been shown to reduce the detection performance of Faster R-CNN (Wang et al. 2022).

5 Relevance to the Present Study

The footwear feature automation can be framed as an object detection task that consists of two main stages: first locating footwear class characteristics or patterns, and classifying the detected regions. However, approaching the problem is more than adapting the network architecture.

5.1 Challenge

Object detection requires precisely annotated training data so that the network can learn both spatial and categorical information. Previous studies have reported that degraded annotation quality can significantly reduce detection performance (Murrugarra-Llerena, Kirsten, and Jung 2022; Wang et al. 2022). Therefore, annotation quality is an important factor in network performance.

Unlike general-purpose pretraining datasets, where object classes are typically clearly defined, footwear features can be more complex and ambiguous. For example, Figure 11a shows a feature coded as a "circle" in the forensic footwear tool SoleMate. However, the feature itself consists of four individual arcs. It is therefore open to discuss whether this feature should be annotated as one bounding region labeled as a circle or as four separate regions labeled as arcs.

In addition to this ambiguity in grouping individual elements, uncertainty can also arise from unclear boundaries between feature categories. For example, Figure 11b and Figure 11c are

coded as “oval” and “bar with rounded ends,” respectively, and belong to two different categories. However, the perceptual boundary between these categories may become unclear as the shape is progressively elongated or compressed.

It is also difficult to exhaustively represent all possible feature variants when the labeling system is restricted to a limited set of basic categories such as circles, bars, and polygons. For example, Figure 11d represents a variant of a circle. Numerous visually different features may still be coded within the same “circle” category, making it impractical to collect examples of every possible variant for model training.



(a) A broken circle or four arcs



(b) Ellipse



(c) A rounded bar



(d) A circle with a knobbly perimeter

5.2 How Human Makes Judgements

A closer examination of the labeling process suggests that object classes should be defined according to common human visual perception. The trained network is intended to assist footwear examiners, consistency between human interpretation and network labels is important. For this reason, the classification labels should align with human cognition in this study.

5.2.1 Human visual System

Visual information is processed hierarchically along the ventral visual pathway to identify an object. When human sees an object, eye-seeing signals first reach V1, where basic features such as edges, orientations, and contrast are detected. Then the information is passed to V2, where simple features are combined into more complex patterns and contours. The information continues progressing to higher level regions where representations becoming increasingly associated with complex shape and object.

The description is mainly a bottom-up process, where sensory information are integrated to make see an object. However,

5.2.2 Gestalt Principles

5.3 How the Problem Will be Processed

The overall performance of the detection network

The generalization capability of backbones

The

6 Reference

Awang, Adriawan. 1962. "Receptive Fields, Binocular Interaction and Functional Architecture in the Cat's Visual Cortex." *The Journal of Physiology*.

Bodziak, William J. 2017. *Footwear Impression Evidence: Detection, Recovery and Examination*. CRC Press.

Deng, Jia, Wei Dong, Richard Socher, Li-Jia Li, Kai Li, and Li Fei-Fei. 2009. "Imagenet: A Large-Scale Hierarchical Image Database." In *2009 IEEE Conference on Computer Vision and Pattern Recognition*, 248–55. Ieee.

Dosovitskiy, Alexey, Lucas Beyer, Alexander Kolesnikov, Dirk Weissenborn, Xiaohua Zhai, Thomas Unterthiner, Mostafa Dehghani, et al. 2020. "An Image Is Worth 16x16 Words: Transformers for Image Recognition at Scale." *arXiv Preprint arXiv:2010.11929*.

Fukushima, Kunihiko. 1980. "Neocognitron: A Self-Organizing Neural Network Model for a Mechanism of Pattern Recognition Unaffected by Shift in Position." *Biological Cybernetics* 36 (4): 193–202.

Girshick, Ross, Jeff Donahue, Trevor Darrell, and Jitendra Malik. 2014. "Rich Feature Hierarchies for Accurate Object Detection and Semantic Segmentation." In *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition*, 580–87.

He, Kaiming, and Jian Sun. 2015. "Convolutional Neural Networks at Constrained Time Cost." In *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition*, 5353–60.

He, Kaiming, Xiangyu Zhang, Shaoqing Ren, and Jian Sun. 2016. "Deep Residual Learning

for Image Recognition." In *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition*, 770–78.

Hubel, David H, and Torsten N Wiesel. 1965. "Receptive Fields and Functional Architecture in Two Nonstriate Visual Areas (18 and 19) of the Cat." *Journal of Neurophysiology* 28 (2): 229–89.

Krizhevsky, Alex, Ilya Sutskever, and Geoffrey E Hinton. 2012. "Imagenet Classification with Deep Convolutional Neural Networks." *Advances in Neural Information Processing Systems* 25.

Lin, Tsung-Yi, Piotr Dollár, Ross Girshick, Kaiming He, Bharath Hariharan, and Serge Belongie. 2017. "Feature Pyramid Networks for Object Detection." In *2017 IEEE Conference on Computer Vision and Pattern Recognition (CVPR)*, 936–44. Ieee.

Murrugarra-Llerena, Jeffri, Lucas N Kirsten, and Claudio R Jung. 2022. "Can We Trust Bounding Box Annotations for Object Detection?" In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, 4813–22.

Simonyan, Karen, and Andrew Zisserman. 2014. "Very Deep Convolutional Networks for Large-Scale Image Recognition." *arXiv Preprint arXiv:1409.1556*.

Vaswani, Ashish, Noam Shazeer, Niki Parmar, Jakob Uszkoreit, Llion Jones, Aidan N Gomez, Łukasz Kaiser, and Illia Polosukhin. 2017. "Attention Is All You Need." *Advances in Neural Information Processing Systems* 30.

Wang, Shaoru, Jin Gao, Bing Li, and Weiming Hu. 2022. "Narrowing the Gap: Improved Detector Training with Noisy Location Annotations." *IEEE Transactions on Image Processing* 31: 6369–80.

Yosinski, Jason, Jeff Clune, Yoshua Bengio, and Hod Lipson. 2014. "How Transferable Are Features in Deep Neural Networks?" *Advances in Neural Information Processing Systems* 27.